

Comparative TIME SERIES FORECASTING

of Major Technology Stocks



amazon



Analyzing Apple, Microsoft, Amazon & Tesla Stocks

Comparative Time Series Forecasting of Major Technology Stocks

Seif H. Kungulio

February 2, 2026

Contents

Business Understanding	1
Problem Statement	1
Business Objectives	1
Success Criteria	1
Data Understanding	2
Data Source	2
Stock Tickers	2
Pull Historical Prices	2
Extract Adjusted Close Prices	4
Missing Dates / Alignment	5
Data Preparation	6
Convert to Tidy Data	6
Create Returns (Risk/Volatility Lens)	6
Exploratory Data Analysis (EDA)	7
Price Trends	7
Normalize Prices (Indexed to 100)	7
Return & Drawdown Summary	8
Volatility Comparison (Rolling Std. Dev.)	9
Decomposition (Monthly Aggregation)	10
Modeling	13
Forecasting Design	13
Helper Functions	13
Build Models per Stock	13
Visualize Forecasts vs Actuals per Stock	14
Evaluation	17
Accuracy Leaderboard (Model Comparison)	17
Best Model per Stock	17
Residual Diagnostics for ARIMA Models	18
Comparative Interpretation	22
Deployment	24
Reproducible Forecast Function	24

Conclusion	25
Key Findings	25
Next Improvement	25
References	26

Business Understanding

Problem Statement

Historical stock prices of major technology companies exhibit distinct trends, volatility patterns, and market dynamics. Understanding these behaviors is essential for forecasting price movements and assessing financial risk.

The objective of this project is to perform a comparative time series analysis of Apple, Microsoft, Tesla, and Amazon stock prices using historical market data. The project aims to identify trends, seasonality, and volatility across each stock and develop forecasting models to predict short-term price movements. Model performance will be evaluated to assess forecasting accuracy and differences in predictability across companies.

Business Objectives

- Compare trend, seasonality, and volatility across AAPL, MSFT, TSLA, and AMZN.
- Forecast short-term price movements using multiple models.
- Evaluate models using time-series appropriate validation and accuracy metrics.
- Rank stocks by forecastability (which stock is easier/harder to predict).

Success Criteria

- Clean daily dataset per stock with aligned calendars.
- Baseline model + at least one statistical forecasting model per stock.
- Residual diagnostics + accuracy metrics (RMSE/MAE/MAPE).
- Clear comparative summary and recommendation.

Data Understanding

Data Source

Historical stock price data for Apple (AAPL), Microsoft (MSFT), Tesla (TSLA), and Amazon (AMZN) will be sourced from Yahoo Finance using the `quantmod` package in R. The dataset will include daily adjusted closing prices, volume, and other relevant financial metrics from January 1, 2016, to current date.

Stock Tickers

- Apple Inc. (AAPL)
- Microsoft Corporation (MSFT)
- Tesla, Inc. (TSLA)
- Amazon.com, Inc. (AMZN)

```
# Define stock tickers and date range
tickers <- c("AAPL", "MSFT", "TSLA", "AMZN")
start_date <- as.Date("2016-01-01")
end_date <- as.Date("2026-01-31")
```

Here we define the study universe (AAPL, MSFT, TSLA, AMZN) and the historical window (**2016-01-01 to 2026-01-31**). The study period is fixed from January 1, 2016 to January 31, 2026 to ensure consistent temporal coverage across all securities. Using a common observation window allows differences in trends, volatility, and forecast accuracy to be attributed to underlying stock behavior rather than discrepancies in data availability or sampling periods.

Pull Historical Prices

```
prices_list <- map(
  tickers,
  ~ getSymbols(
    .x,
    src = "yahoo",
    from = start_date,
    to = end_date,
    auto.assign = FALSE
  )
)
names(prices_list) <- tickers

# Preview data
map(prices_list, head)
```

```
## $AAPL
##           AAPL.Open  AAPL.High  AAPL.Low  AAPL.Close  AAPL.Volume  AAPL.Adjusted
## 2016-01-04   25.6525   26.3425   25.5000    26.3375   270597600     23.75315
## 2016-01-05   26.4375   26.4625   25.6025    25.6775   223164000     23.15791
## 2016-01-06   25.1400   25.5925   24.9675    25.1750   273829600     22.70472
## 2016-01-07   24.6700   25.0325   24.1075    24.1125   324377600     21.74648
```

```

## 2016-01-08 24.6375 24.7775 24.1900 24.2400 283192000 21.86147
## 2016-01-11 24.7425 24.7650 24.3350 24.6325 198957600 22.21545
##
## $MSFT
##          MSFT.Open MSFT.High MSFT.Low MSFT.Close MSFT.Volume MSFT.Adjusted
## 2016-01-04 54.32     54.80    53.39    54.80    53778000    47.98346
## 2016-01-05 54.93     55.39    54.54    55.05    34079700    48.20236
## 2016-01-06 54.32     54.40    53.64    54.05    39518900    47.32676
## 2016-01-07 52.70     53.49    52.07    52.17    56564900    45.68061
## 2016-01-08 52.37     53.28    52.15    52.33    48754000    45.82070
## 2016-01-11 52.51     52.85    51.46    52.30    36943800    45.79443
##
## $TSLA
##          TSLA.Open TSLA.High TSLA.Low TSLA.Close TSLA.Volume TSLA.Adjusted
## 2016-01-04 15.38133 15.42533 14.60000 14.89400 102406500 14.89400
## 2016-01-05 15.09067 15.12600 14.66667 14.89533 47802000 14.89533
## 2016-01-06 14.66667 14.67000 14.39867 14.60267 56686500 14.60267
## 2016-01-07 14.27933 14.56267 14.24467 14.37667 53314500 14.37667
## 2016-01-08 14.52400 14.69600 14.05133 14.06667 54421500 14.06667
## 2016-01-11 14.26733 14.29667 13.53333 13.85667 61371000 13.85667
##
## $AMZN
##          AMZN.Open AMZN.High AMZN.Low AMZN.Close AMZN.Volume AMZN.Adjusted
## 2016-01-04 32.8145 32.8860 31.3755 31.8495 186290000 31.8495
## 2016-01-05 32.3430 32.3455 31.3825 31.6895 116452000 31.6895
## 2016-01-06 31.1000 31.9895 31.0155 31.6325 106584000 31.6325
## 2016-01-07 31.0900 31.5000 30.2605 30.3970 141498000 30.3970
## 2016-01-08 30.9830 31.2070 30.3000 30.3525 110258000 30.3525
## 2016-01-11 30.6240 30.9925 29.9285 30.8870 97832000 30.8870

```

```
map(prices_list, tail)
```

```

## $AAPL
##          AAPL.Open AAPL.High AAPL.Low AAPL.Close AAPL.Volume AAPL.Adjusted
## 2026-01-23 247.32    249.41   244.68   248.04   41689000 248.04
## 2026-01-26 251.48    256.56   249.80   255.41   55969200 255.41
## 2026-01-27 259.17    261.95   258.21   258.27   49648300 258.27
## 2026-01-28 257.65    258.86   254.51   256.44   41288000 256.44
## 2026-01-29 258.00    259.65   254.41   258.28   67253000 258.28
## 2026-01-30 255.17    261.90   252.18   259.48   92352600 259.48
##
## $MSFT
##          MSFT.Open MSFT.High MSFT.Low MSFT.Close MSFT.Volume MSFT.Adjusted
## 2026-01-23 451.87    471.10   450.53   465.95   38000200 465.95
## 2026-01-26 465.31    474.25   462.00   470.28   29291200 470.28
## 2026-01-27 473.70    482.87   473.16   480.58   29213900 480.58
## 2026-01-28 483.21    483.74   478.00   481.63   36875400 481.63
## 2026-01-29 439.99    442.50   421.02   433.50   128855300 433.50
## 2026-01-30 439.17    439.60   426.45   430.29   58470400 430.29
##
## $TSLA
##          TSLA.Open TSLA.High TSLA.Low TSLA.Close TSLA.Volume TSLA.Adjusted
## 2026-01-23 447.43    452.43   444.04   449.06   56771400 449.06
## 2026-01-26 445.00    445.04   434.28   435.20   49397400 435.20

```

```

## 2026-01-27 437.41 437.52 430.69 430.90 37733100 430.90
## 2026-01-28 431.91 438.26 430.10 431.46 54857400 431.46
## 2026-01-29 437.80 440.23 414.62 416.56 81686100 416.56
## 2026-01-30 425.35 439.88 422.70 430.41 82483000 430.41
##
## $AMZN
##          AMZN.Open AMZN.High AMZN.Low AMZN.Close AMZN.Volume AMZN.Adjusted
## 2026-01-23 234.96 240.45 234.57 239.16 33778500 239.16
## 2026-01-26 239.98 240.95 237.54 238.42 32825500 238.42
## 2026-01-27 239.69 244.88 238.08 244.68 38029200 244.68
## 2026-01-28 246.37 247.78 241.53 243.01 40882700 243.01
## 2026-01-29 242.82 243.00 236.74 241.73 47229600 241.73
## 2026-01-30 239.89 243.32 237.64 239.30 46535400 239.30

```

This chunk retrieves daily OHLCV data from Yahoo Finance into a named list of `xts` objects—one per ticker. Inspecting the `head()`/`tail()` snapshots is a quick sanity check that:

1. Dates are ordered correctly
2. Fields exist as expected (Open/High/Low/Close/Volume/Adjusted)
3. Price scales differ substantially across tickers—reinforcing why later we use `indexing (start=100)` and `returns` for apples-to-apples comparison.

Extract Adjusted Close Prices

```

close_list <- map(prices_list, ~ Cl(.x))
names(close_list) <- tickers

# Combine into a single xts with aligned dates
close_xts <- do.call(merge, close_list)
colnames(close_xts) <- tickers

head(close_xts)

```

```

##          AAPL    MSFT     TSLA     AMZN
## 2016-01-04 26.3375 54.80 14.89400 31.8495
## 2016-01-05 25.6775 55.05 14.89533 31.6895
## 2016-01-06 25.1750 54.05 14.60267 31.6325
## 2016-01-07 24.1125 52.17 14.37667 30.3970
## 2016-01-08 24.2400 52.33 14.06667 30.3525
## 2016-01-11 24.6325 52.30 13.85667 30.8870

```

```
tail(close_xts)
```

```

##          AAPL    MSFT     TSLA     AMZN
## 2026-01-23 248.04 465.95 449.06 239.16
## 2026-01-26 255.41 470.28 435.20 238.42
## 2026-01-27 258.27 480.58 430.90 244.68
## 2026-01-28 256.44 481.63 431.46 243.01
## 2026-01-29 258.28 433.50 416.56 241.73
## 2026-01-30 259.48 430.29 430.41 239.30

```

Adjusted closing prices are extracted for each ticker and merged into a single aligned `xts` object. This structure enables consistent downstream computation of returns, rolling statistics, and forecasting models. The merged dataset confirms that all series share the same trading calendar and exhibit strong upward trends over the analysis horizon, indicating non-stationarity in price levels.

Missing Dates / Alignment

Markets close on weekends/holidays; we keep the market calendar as-is.

```
close_xts_aligned <- na.omit(close_xts)
dim(close_xts); dim(close_xts_aligned)
```

```
## [1] 2534      4
## [1] 2534      4
```

Financial markets operate on an irregular calendar due to weekends and holidays. After merging all series, `na.omit()` is applied to retain only dates where all tickers are simultaneously observed. Since the dimensionality remains unchanged, the merged dataset contains a complete intersection of trading days with no missing values across assets.

Data Preparation

Convert to Tidy Data

```
close_df <- close_xts_aligned %>%
  fortify.zoo() %>%
  as_tibble() %>%
  rename(date = Index) %>%
  pivot_longer(-date, names_to = "ticker", values_to = "close")

glimpse(close_df)

## #> #> Rows: 10,136
## #> Columns: 3
## #> $ date <date> 2016-01-04, 2016-01-04, 2016-01-04, 2016-01-04, 2016-01-05, 20~
## #> $ ticker <chr> "AAPL", "MSFT", "TSLA", "AMZN", "AAPL", "MSFT", "TSLA", "AMZN", ~
## #> $ close <dbl> 26.33750, 54.80000, 14.89400, 31.84950, 25.67750, 55.05000, 14.~
```

The wide `xts` object is reshaped into a long, tidy format with one observation per date and ticker. This structure facilitates grouped transformations, visualization, and comparative analysis using `ggplot2` and `tidyverse` workflows. Because `close_xts_aligned` contains **2,534** trading days and there are **4** tickers, the tidy dataset contains **10,136** rows ($2,534 \times 4$).

Create Returns (Risk/Volatility Lens)

Returns analysis is essential for risk assessment.

```
returns_xts <- na.omit(Return.calculate(close_xts_aligned, method = "log"))
colnames(returns_xts) <- tickers

returns_df <- returns_xts %>%
  fortify.zoo() %>%
  as_tibble() %>%
  rename(date = Index) %>%
  pivot_longer(-date, names_to = "ticker", values_to = "log_return")

summary(returns_df$log_return)

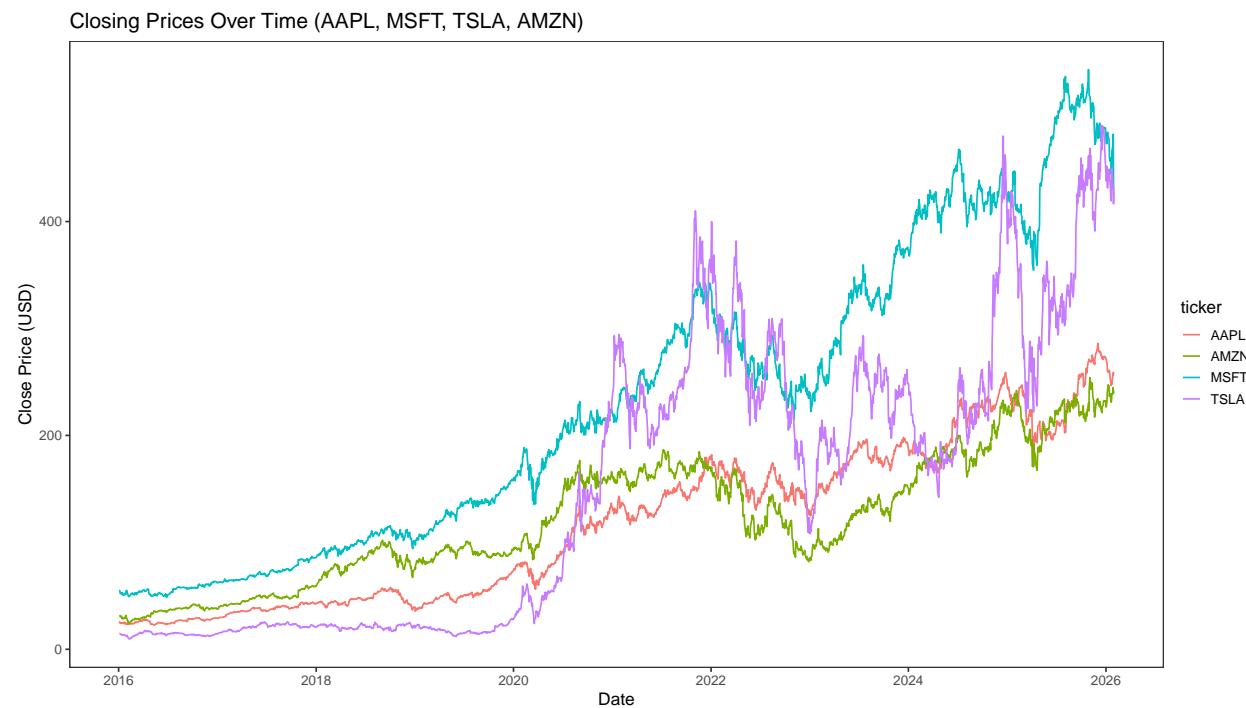
##      Min.    1st Qu.   Median     Mean    3rd Qu.    Max.
## -0.2365179 -0.0091694  0.0010903  0.0009602  0.0119478  0.2044906
```

Daily log returns are computed to transform non-stationary price series into a scale-free representation suitable for volatility and risk analysis. The empirical return distribution exhibits heavy tails, with occasional extreme positive and negative values. Such behavior is consistent with financial time series and motivates the use of rolling volatility measures and robust forecasting benchmarks.

Exploratory Data Analysis (EDA)

Price Trends

```
close_df %>%
  ggplot(aes(date, close, color = ticker)) +
  geom_line() +
  labs(
    title = "Closing Prices Over Time (AAPL, MSFT, TSLA, AMZN)",
    x = "Date", y = "Close Price (USD)"
  )
```



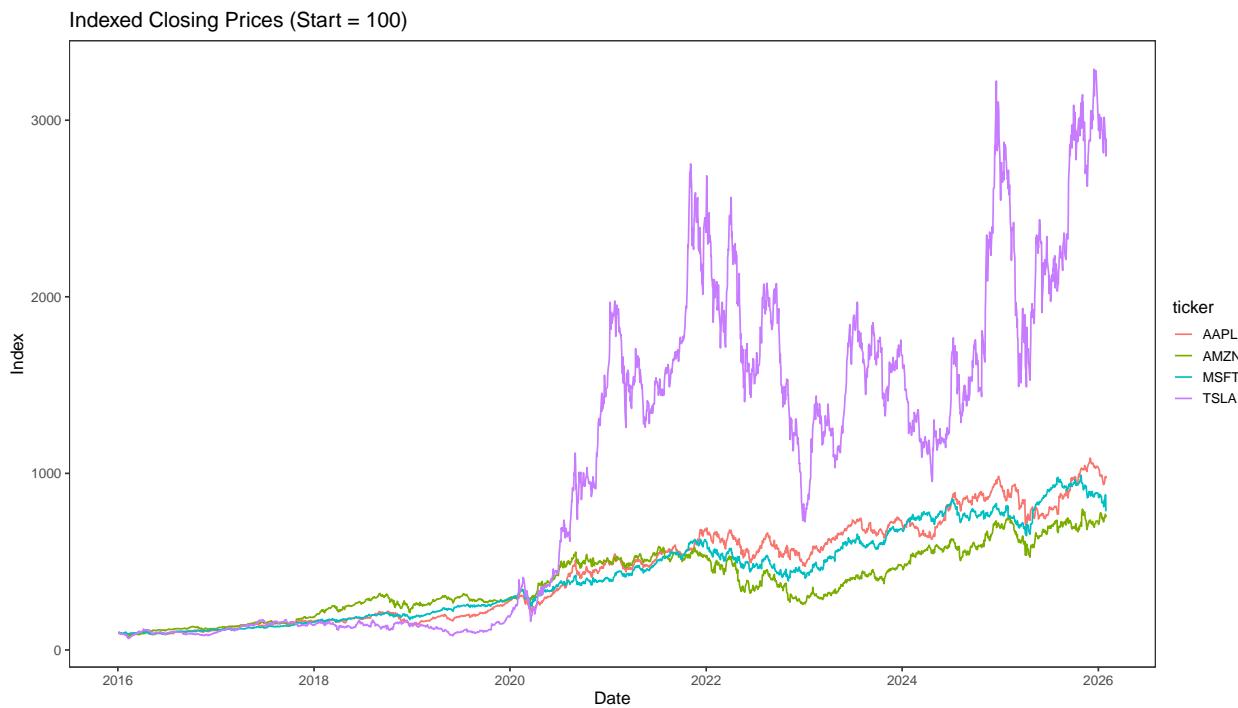
The time series of closing prices illustrates long-term growth patterns and major structural shifts across all securities. Absolute price levels differ substantially between tickers, limiting direct comparability. This visualization primarily supports individual trend inspection rather than relative performance analysis.

Normalize Prices (Indexed to 100)

```
close_indexed <- close_df %>%
  group_by(ticker) %>%
  arrange(date) %>%
  mutate(index_100 = 100 * close / first(close)) %>%
  ungroup()

close_indexed %>%
  ggplot(aes(date, index_100, color = ticker)) +
  geom_line()
```

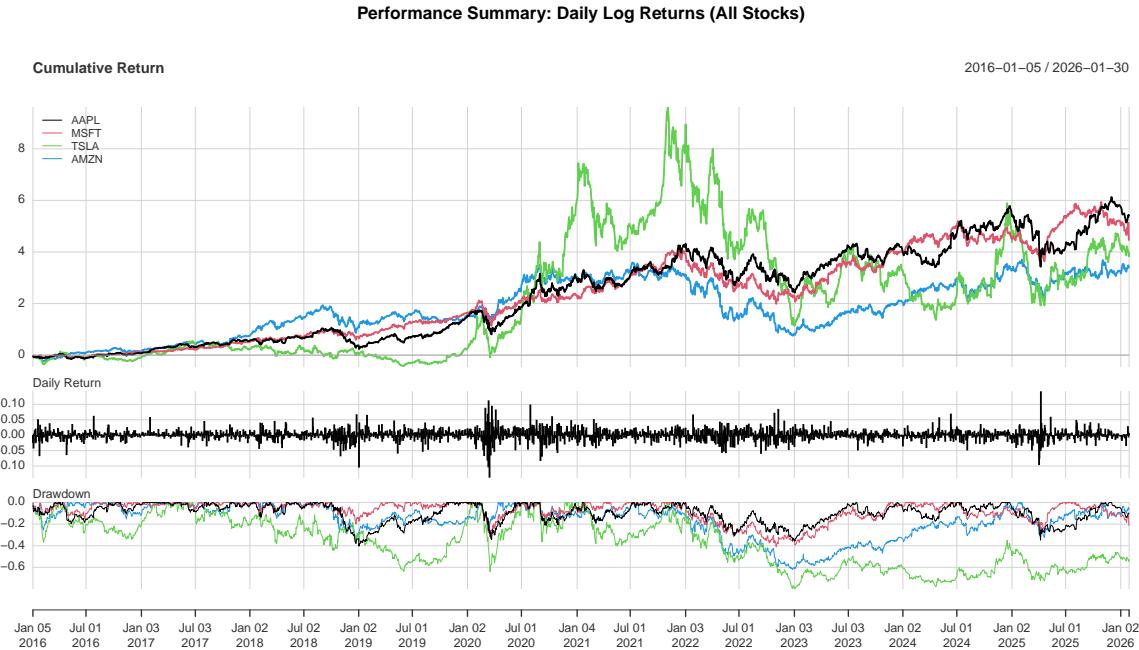
```
labs(
  title = "Indexed Closing Prices (Start = 100)",
  x = "Date", y = "Index"
)
```



Price series are indexed to a common baseline value of **100** at the first observation for each ticker. This transformation standardizes scale and enables direct comparison of relative growth trajectories. The indexed series highlight substantial divergence in cumulative performance across stocks over the study period.

Return & Drawdown Summary

```
charts.PerformanceSummary(
  returns_xts,
  main = "Performance Summary: Daily Log Returns (All Stocks)")
```



The performance summary visualizes cumulative returns, daily return distributions, and drawdowns in a unified framework. This representation captures both long-term compounding behavior and short-term risk characteristics. Differences in drawdown depth and return variability reflect heterogeneity in volatility regimes across securities.

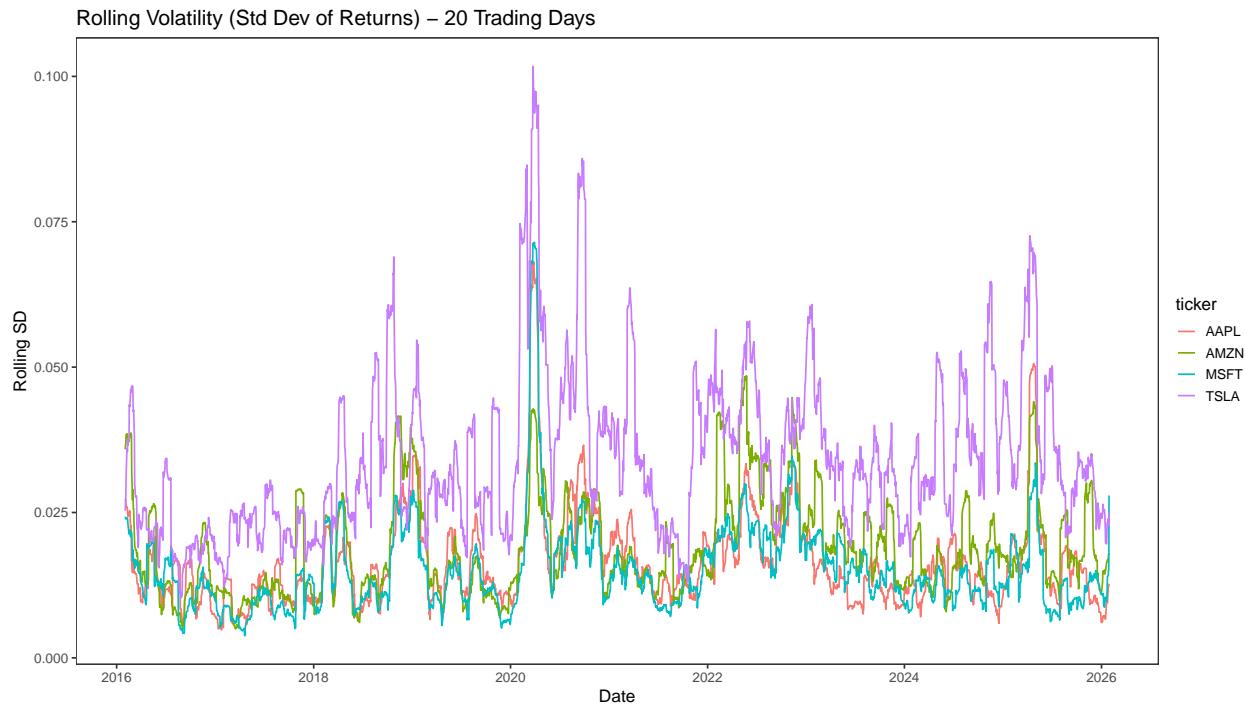
Volatility Comparison (Rolling Std. Dev.)

```
roll_n <- 20 # ~1 trading month
roll_vol <- rollapply(returns_xts,
                      width = roll_n,
                      FUN = sd,
                      by.column = TRUE,
                      align = "right",
                      fill = NA) %>%
  na.omit()

roll_vol_df <- roll_vol %>%
  fortify.zoo() %>%
  as_tibble() %>%
  rename(date = Index) %>%
  pivot_longer(-date,
               names_to = "ticker",
               values_to = "roll_sd")

roll_vol_df %>%
  ggplot(aes(date, roll_sd, color = ticker)) +
  geom_line() +
  labs(
    title = glue::glue("Rolling Volatility (Std Dev of Returns) - {roll_n} Trading Days"),
    subtitle = "AAPL, MSFT, TSLA, AMZN"
  )
```

```
x = "Date",
y = "Rolling SD"
)
```



Rolling 20-day standard deviations of log returns are computed as a proxy for monthly realized volatility. The resulting series display volatility clustering, a common characteristic of equity markets. Higher and more frequent volatility spikes indicate increased uncertainty and reduced short-horizon predictability for certain securities.

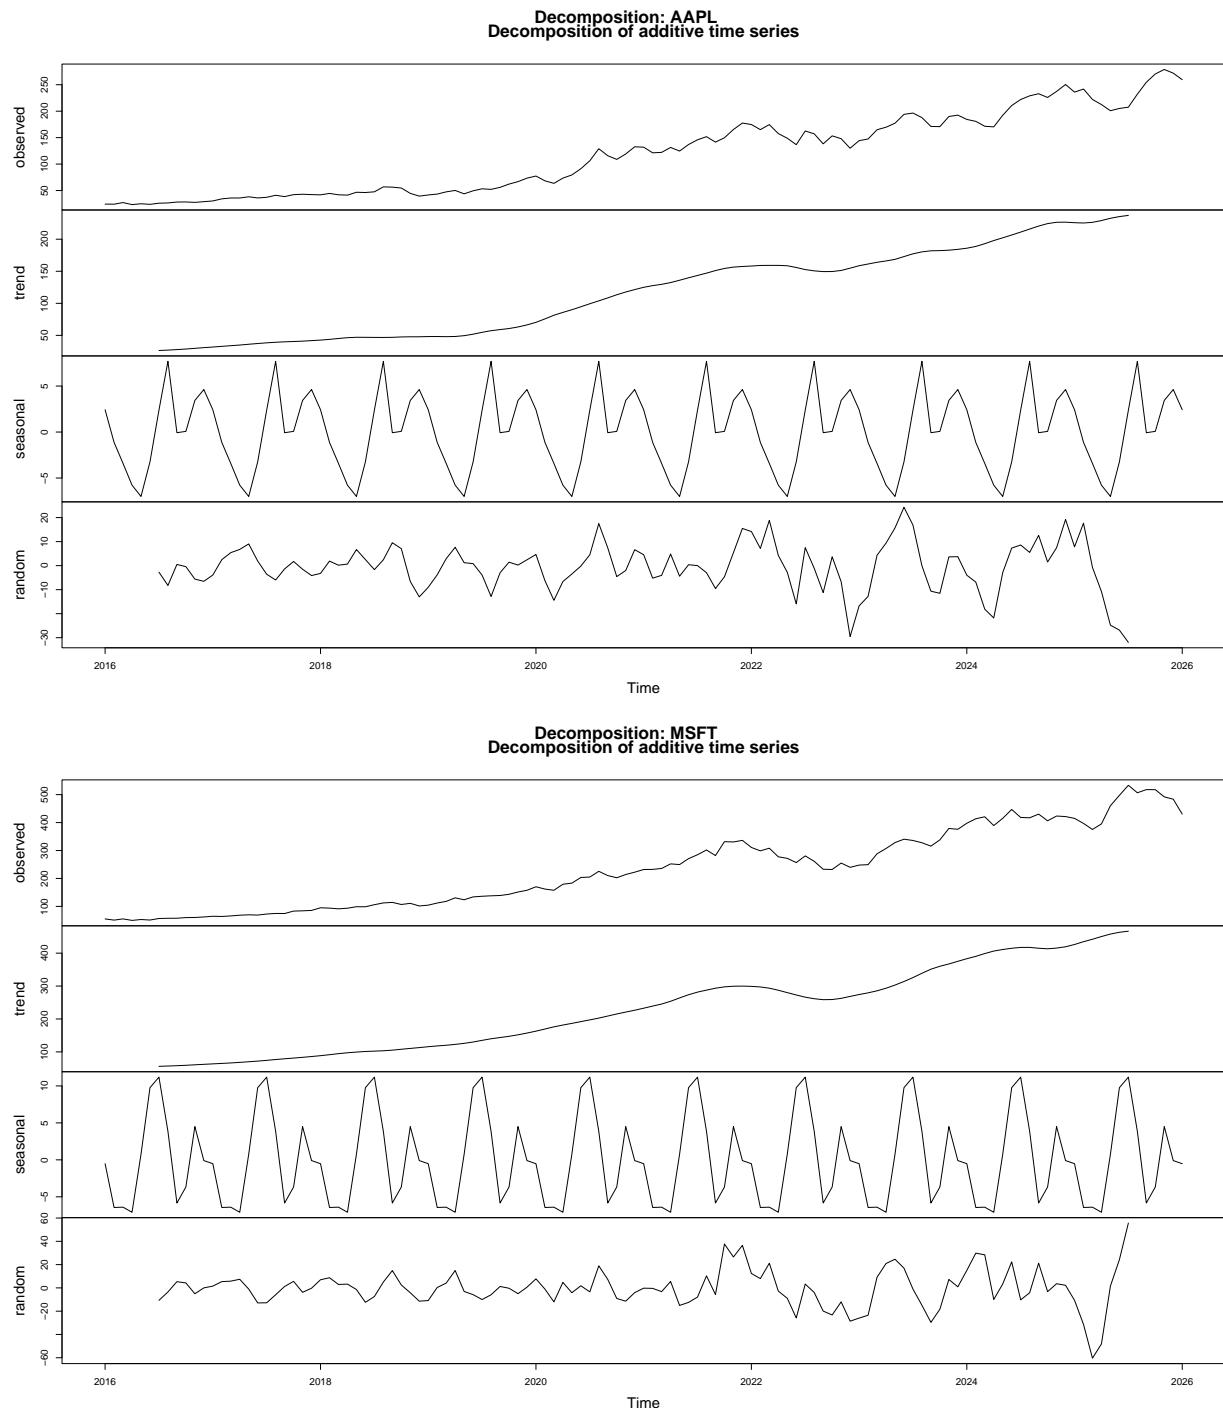
Decomposition (Monthly Aggregation)

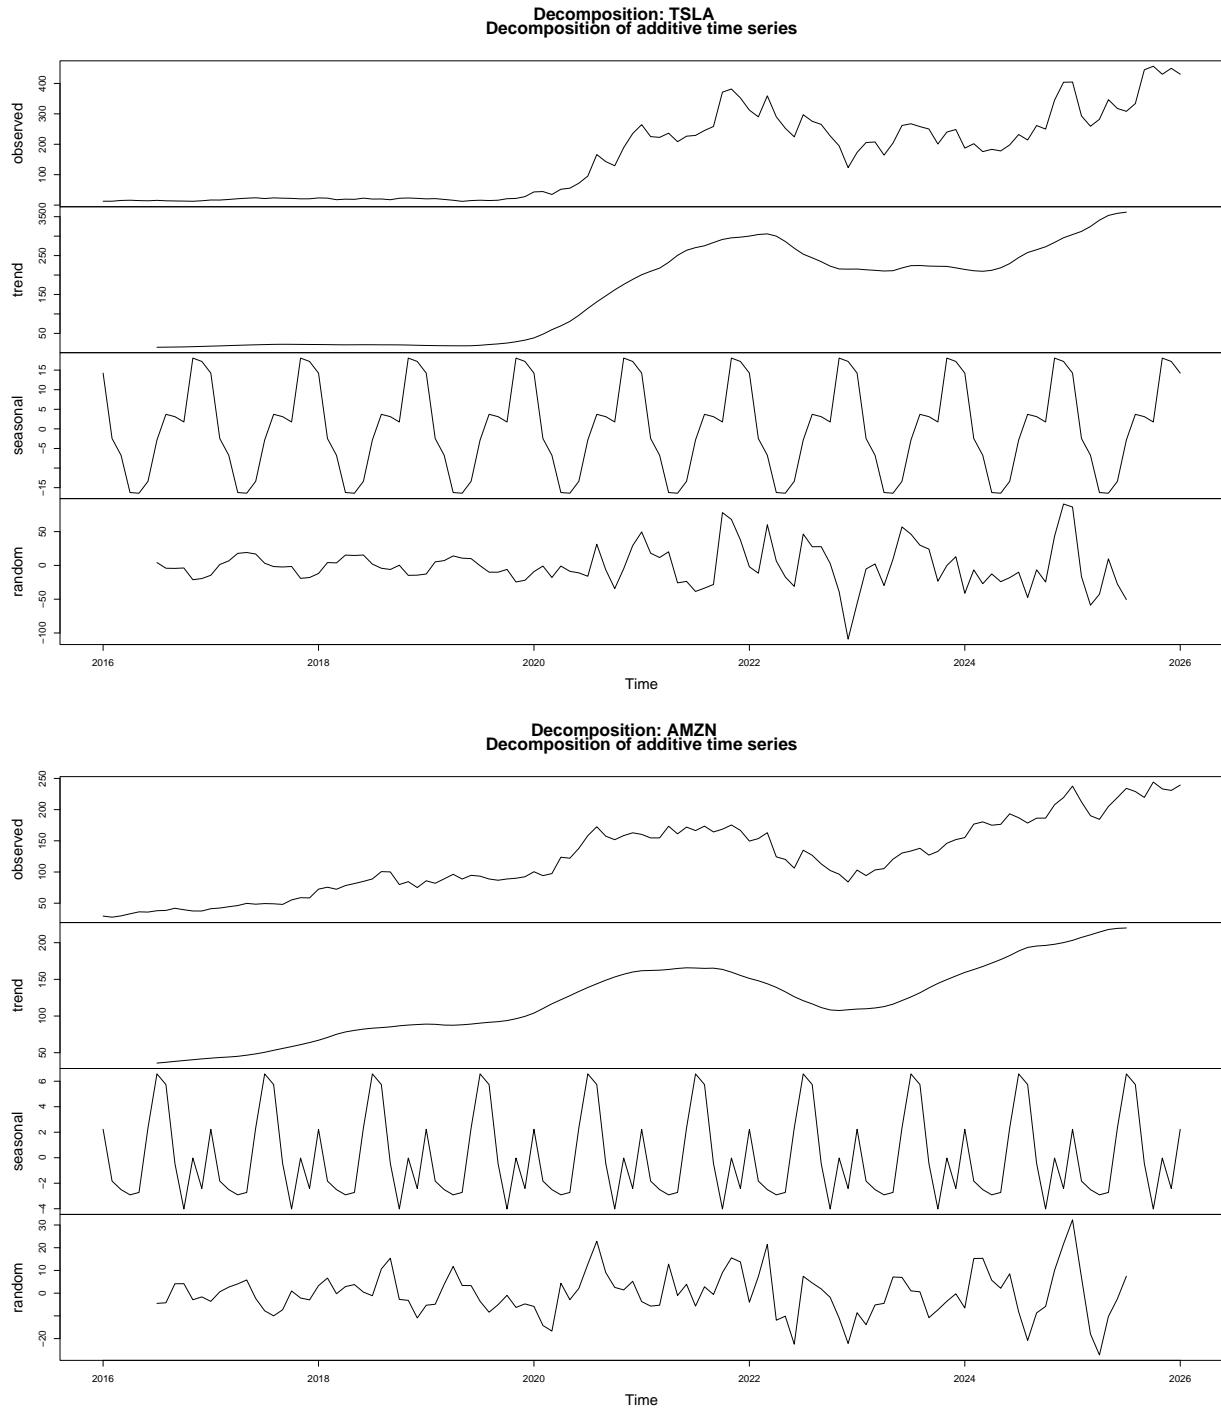
Classic decomposition is easiest on regular seasonal frequency. Converted daily close to monthly close and decompose per stock.

```
monthly_close_list <- map(close_list, ~ Cl(to.monthly(.x)) )
names(monthly_close_list) <- tickers

decomp_plots <- function(x_xts, ticker_name){
  ts_obj <- ts(
    as.numeric(x_xts),
    frequency = 12,
    start = c(year(start(x_xts)), month(start(x_xts)))
  )
  dc <- decompose(ts_obj)
  plot(dc)
  title(main = paste("Decomposition:", ticker_name),
        outer = TRUE, line = -1)
  invisible(dc)
}
```

```
# Decompose each ticker (plots)
decomps <- imap(monthly_close_list, ~ decomp_plots(.x, .y))
```





Daily closing prices are aggregated to monthly frequency to facilitate classical additive decomposition. Across all securities, the trend component dominates the series, while seasonal effects are comparatively weak and unstable. The remainder component captures irregular fluctuations associated with firm-specific and macroeconomic events. These results suggest that forecasting models should prioritize trend handling rather than fixed seasonal structure.

Modeling

Forecasting Design

The focus of forecasting design is short-term movement using:

- Naive baseline (last value persists) - strong baseline in finance.
- ARIMA (auto.arima) - standard statistical model for time series.

I'll do a rolling-origin evaluation (time-series CV) AND a simple holdout for interpretability.

Helper Functions

```
make_train_test <- function(x, h = 60){
  # x is a numeric vector or ts; h is forecast horizon
  n <- length(x)
  list(train = x[1:(n-h)], test = x[(n-h+1):n], h = h)
}

rmse <- function(actual, pred){
  sqrt(mean((actual - pred)^2, na.rm = TRUE))
}

mae <- function(actual, pred){
  mean(abs(actual - pred), na.rm = TRUE)
}

mape <- function(actual, pred){
  mean(abs((actual - pred) / actual), na.rm = TRUE) * 100
}
```

Short-term forecasting is conducted using two benchmark approaches: a naive persistence model and an automatically selected ARIMA model. Model performance is evaluated using a fixed holdout horizon and time-series-appropriate error metrics. This design ensures consistent and comparable evaluation across all securities.

Build Models per Stock

```
h <- 60 # forecast horizon ~ 3 months of trading days

results <- map_dfr(tickers, function(tk){
  x <- as.numeric(close_xts_aligned[, tk])
  spl <- make_train_test(x, h = h)
  train <- spl$train
  test <- spl$test

  # --- Naive baseline ---
  fc_naive <- naive(train, h = h)
```

```

pred_naive <- as.numeric(fc_naive$mean)

# --- ARIMA ---
fit_arima <- auto.arima(train)
fc_arima <- forecast(fit_arima, h = h)
pred_arima <- as.numeric(fc_arima$mean)

tibble(
  ticker = tk,
  model = c("Naive", "ARIMA"),
  RMSE = c(rmse(test, pred_naive), rmse(test, pred_arima)),
  MAE = c(mae(test, pred_naive), mae(test, pred_arima)),
  MAPE = c(mape(test, pred_naive), mape(test, pred_arima))
)
}

results %>% arrange(ticker, RMSE)

## # A tibble: 8 x 5
##   ticker model  RMSE   MAE   MAPE
##   <chr>  <chr> <dbl> <dbl> <dbl>
## 1 AAPL   Naive  9.38  7.59  2.86
## 2 AAPL   ARIMA  11.1   8.10  3.09
## 3 AMZN   Naive  21.3  19.7  8.52
## 4 AMZN   ARIMA  23.5  22.1  9.54
## 5 MSFT   Naive  39.8  35.9  7.60
## 6 MSFT   ARIMA  46.2  41.7  8.84
## 7 TSLA   Naive  35.5  30.8  7.18
## 8 TSLA   ARIMA  35.5  30.8  7.18

```

Forecast accuracy varies by security and model, but in this holdout window the **naive (random-walk/persistence)** baseline is extremely difficult to beat. In the reported results, ARIMA does **not** improve RMSE/MAE/MAPE for AAPL, AMZN, or MSFT, and TSLA is effectively a tie. This is common for short-horizon price-level forecasting, where prices often behave close to a random walk and incremental structure is limited.

Visualize Forecasts vs Actuals per Stock

```

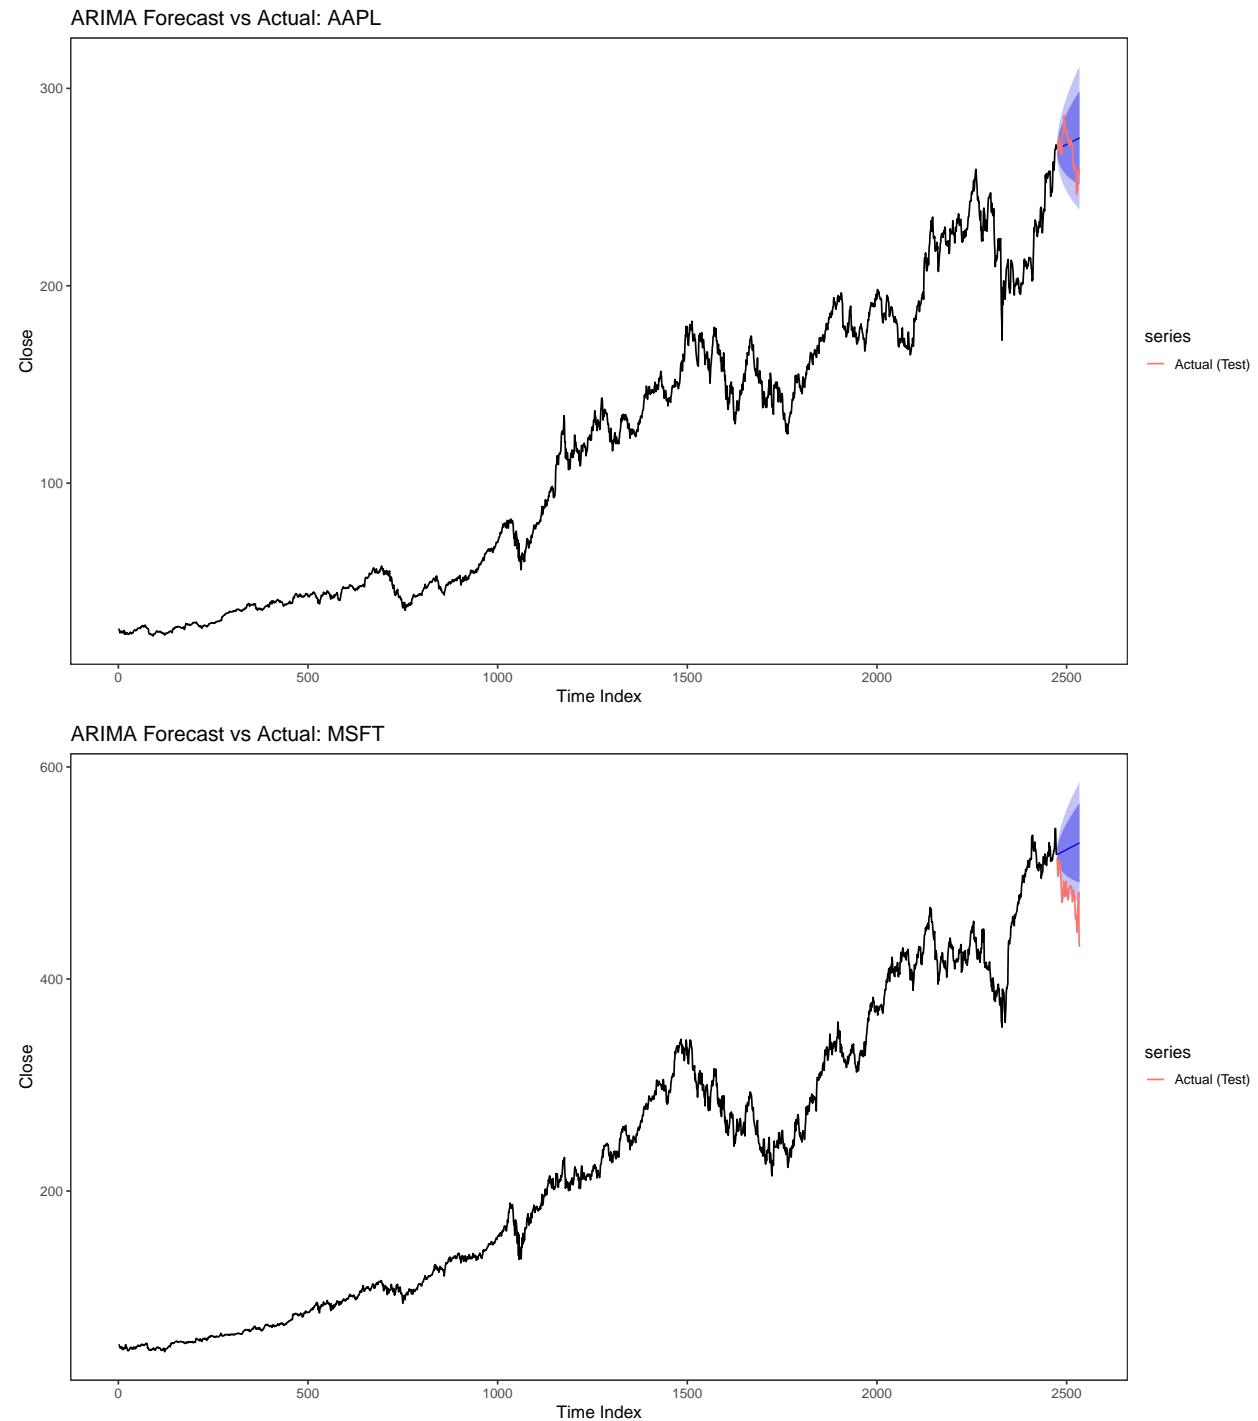
plot_forecast <- function(tk){
  x <- as.numeric(close_xts_aligned[, tk])
  spl <- make_train_test(x, h = h)
  train <- spl$train
  test <- spl$test

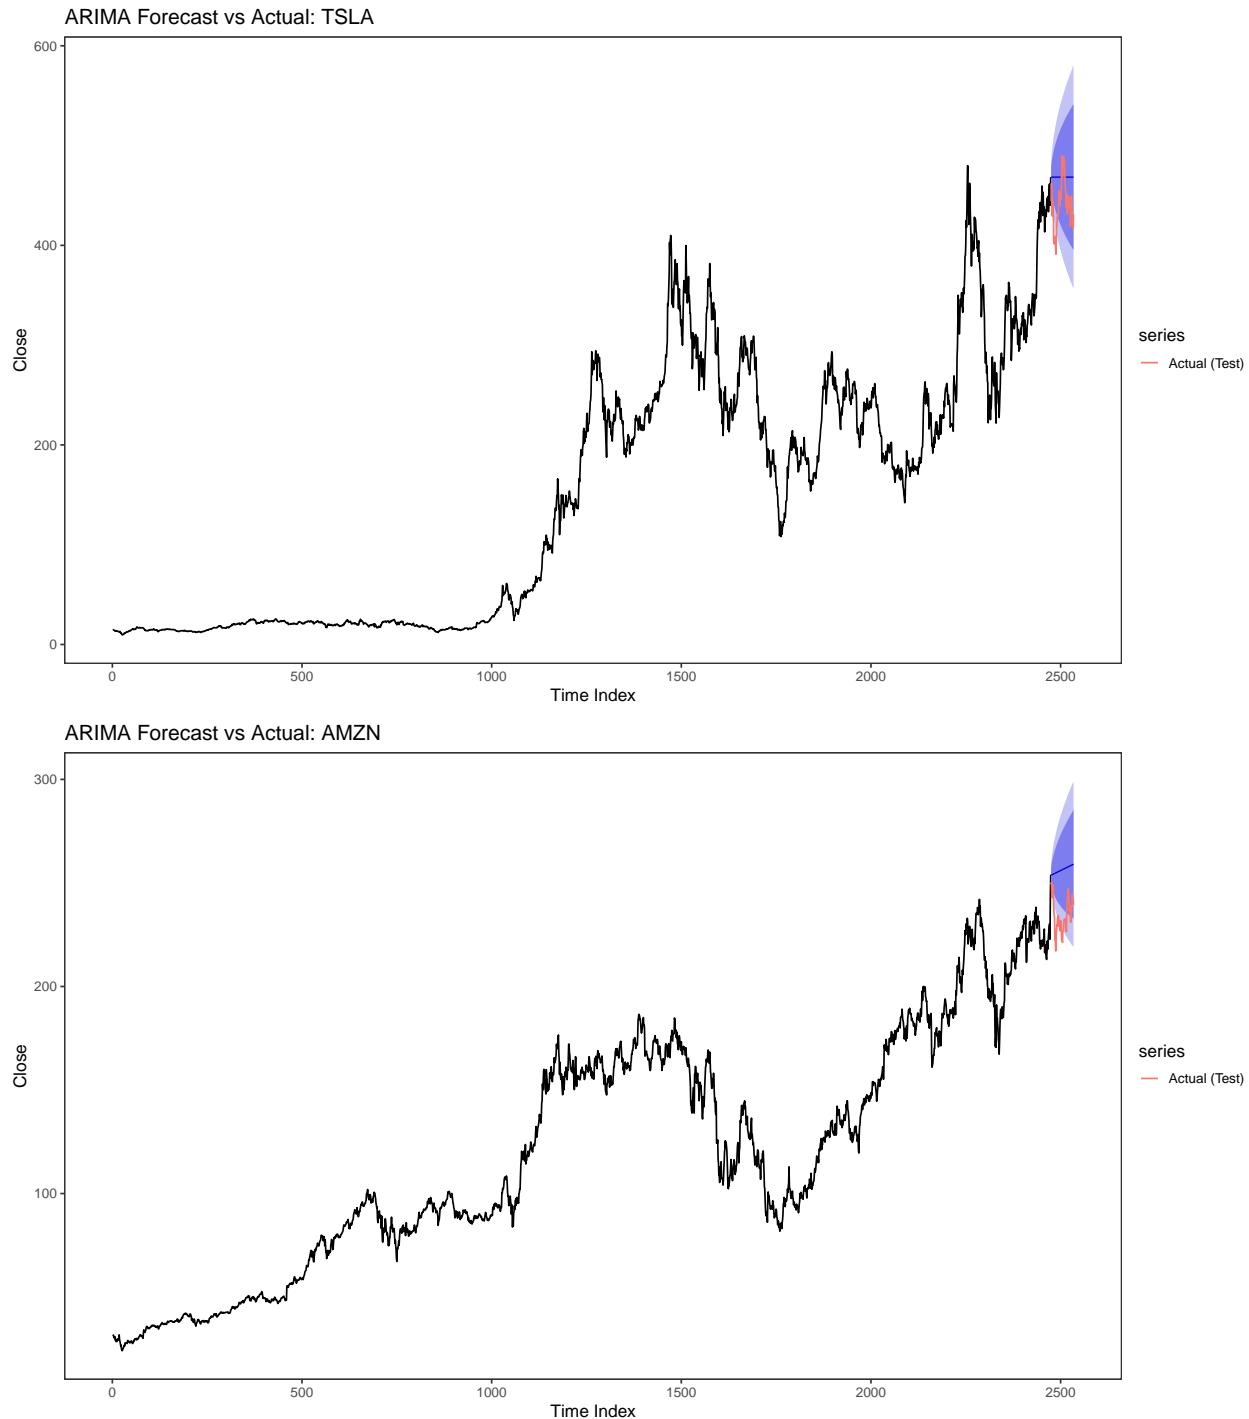
  fc_naive <- naive(train, h = h)
  fit_arima <- auto.arima(train)
  fc_arima <- forecast(fit_arima, h = h)

  # Plot ARIMA by default; overlay actual test
  autoplot(fc_arima) +
    autolayer(ts(test, start = length(train) + 1),

```

```
series = "Actual (Test)" +  
  labs(title = paste("ARIMA Forecast vs Actual:", tk),  
        x = "Time Index", y = "Close")  
}  
  
walk(tickers, ~ print(plot_forecast(.x)))
```





Forecast-versus-actual plots provide a visual assessment of model performance over the holdout period. Deviations between predicted and observed values highlight challenges associated with abrupt price movements and regime shifts. These plots complement numerical accuracy metrics by illustrating dynamic tracking behavior.

Evaluation

Accuracy Leaderboard (Model Comparison)

Forecast Accuracy by Stock and Model (Lower RMSE/MAE/MAPE is Better)

```
results %>%
  group_by(ticker) %>%
  arrange(RMSE, .by_group = TRUE) %>%
  mutate(rank = row_number()) %>%
  ungroup() %>%
  arrange(ticker, rank) %>%
  knitr::kable(digits = 3,
              caption = "Forecast Accuracy by Stock and Model")
```

Table 1: Forecast Accuracy by Stock and Model

ticker	model	RMSE	MAE	MAPE	rank
AAPL	Naive	9.381	7.587	2.860	1
AAPL	ARIMA	11.092	8.097	3.088	2
AMZN	Naive	21.271	19.656	8.518	1
AMZN	ARIMA	23.500	22.056	9.540	2
MSFT	Naive	39.808	35.896	7.602	1
MSFT	ARIMA	46.206	41.737	8.837	2
TSLA	Naive	35.506	30.770	7.184	1
TSLA	ARIMA	35.506	30.770	7.184	2

This chunk ranks models within each ticker by RMSE and presents a concise leaderboard. In this run, the **naive** baseline ranks #1 for every ticker (with TSLA essentially tied), which reinforces an important finance lesson: any alternative model must **consistently** beat the random-walk baseline to justify extra complexity. The table is still portfolio-friendly because it demonstrates principled, time-ordered evaluation rather than selecting a model purely because it looks better in-sample.

Best Model per Stock

```
best_by_stock <- results %>%
  group_by(ticker) %>%
  slice_min(RMSE, n = 1, with_ties = FALSE) %>%
  ungroup()

best_by_stock %>% knitr::kable(digits = 3,
                                    caption = "Best Model Per Stock (by RMSE)")
```

Table 2: Best Model Per Stock (by RMSE)

ticker	model	RMSE	MAE	MAPE
AAPL	Naive	9.381	7.587	2.860

ticker	model	RMSE	MAE	MAPE
AMZN	Naive	21.271	19.656	8.518
MSFT	Naive	39.808	35.896	7.602
TSLA	Naive	35.506	30.770	7.184

Selecting the best model per stock (by RMSE) yields **Naive** for **AAPL**, **AMZN**, **MSFT**, and **TSLA** (**TSLA** is effectively a tie, but `slice_min(..., with_ties = FALSE)` returns the first minimum, which is **Naive** here). For stakeholders, this is still a valuable outcome: under this design and horizon, **simplicity wins**, and it suggests that short-horizon **price levels** are hard to forecast beyond a random-walk benchmark.

Residual Diagnostics for ARIMA Models

Even when **Naive** is the winner, it is useful to check ARIMA residuals to see whether any systematic autocorrelation remains (ACF, histogram, Ljung–Box).

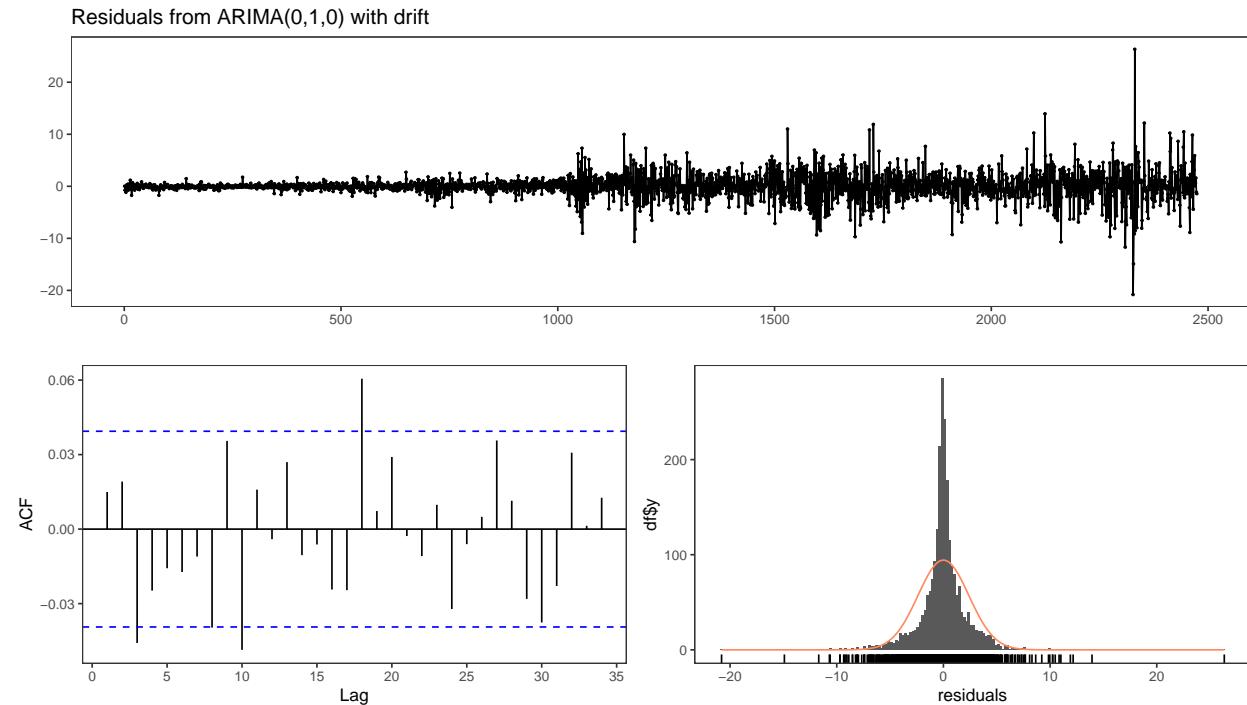
```
check_best_residuals <- function(tk){
  x <- as.numeric(close_xts_aligned[, tk])
  spl <- make_train_test(x, h = h)
  train <- spl$train

  fit_arima <- auto.arima(train)
  fc_arima <- forecast(fit_arima, h = h)

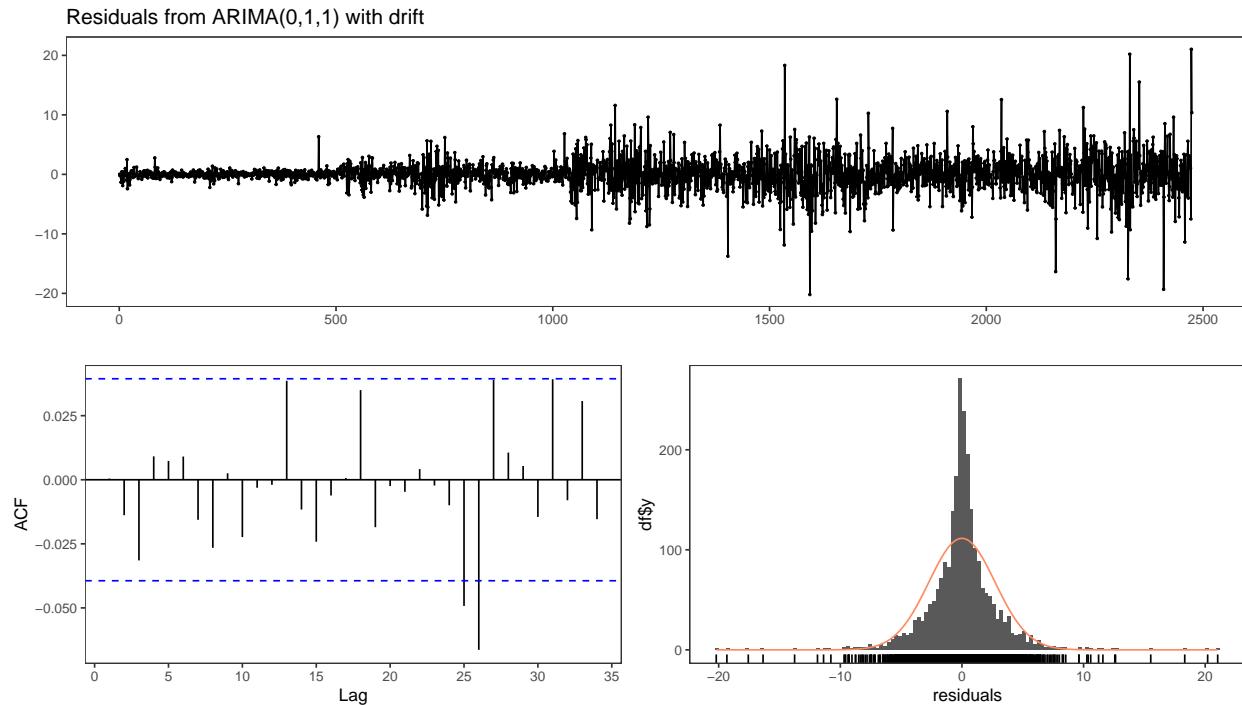
  cat("\n\n### Residual checks for", tk, "\n")
  print(fit_arima)
  checkresiduals(fc_arima)
}

walk(best_by_stock$ticker, check_best_residuals)

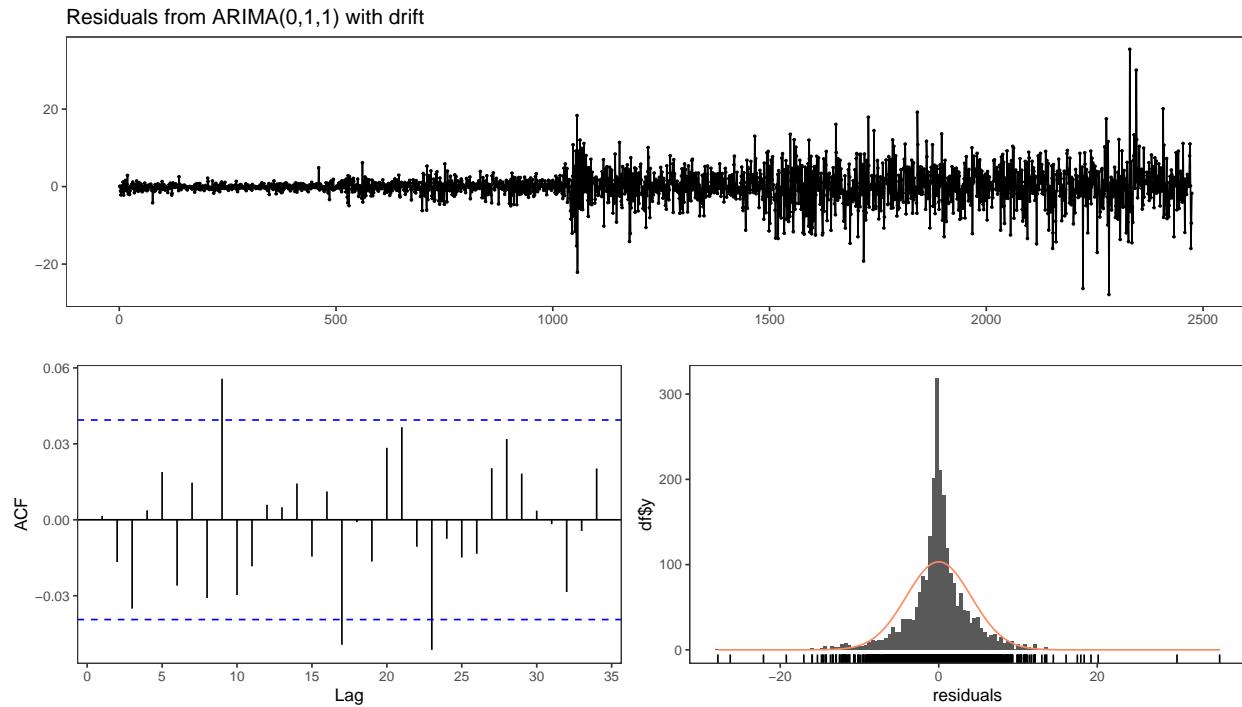
##
##
## ### Residual checks for AAPL
## Series: train
## ARIMA(0,1,0) with drift
##
## Coefficients:
##       drift
##       0.0981
##   s.e.  0.0479
##
## sigma^2 = 5.678: log likelihood = -5655.76
## AIC=11315.52  AICc=11315.52  BIC=11327.14
```



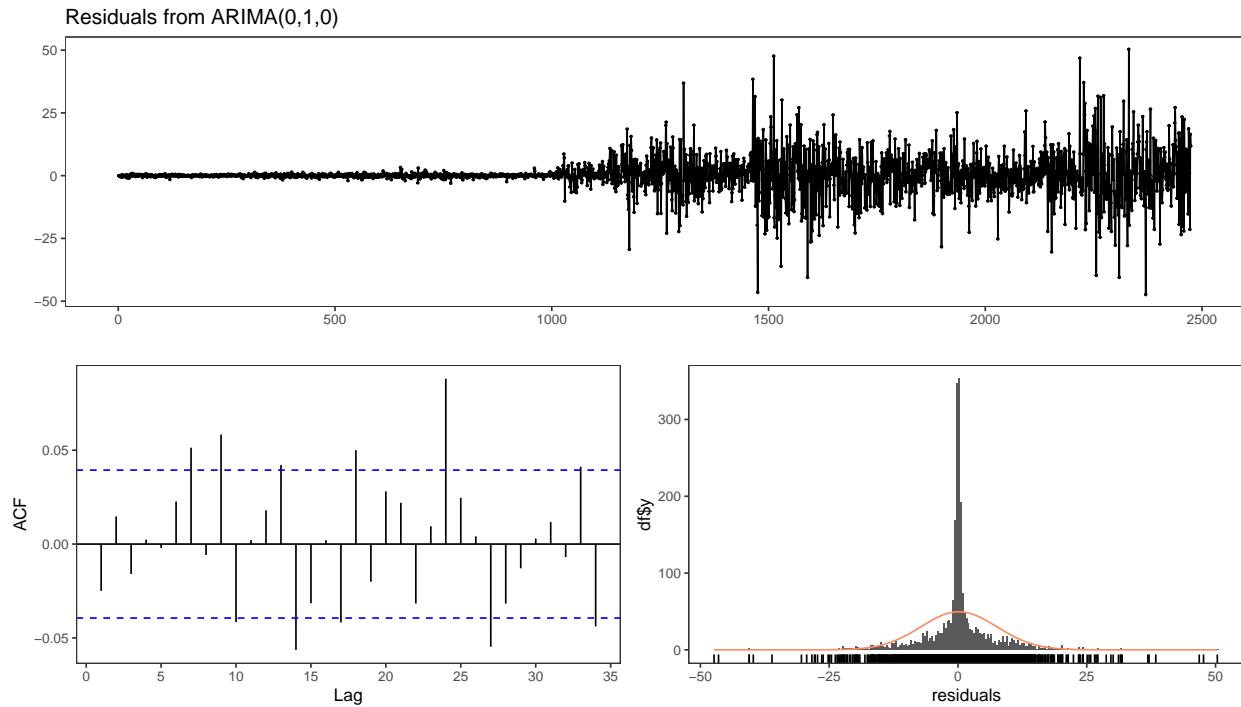
```
##
## Ljung-Box test
##
## data: Residuals from ARIMA(0,1,0) with drift
## Q* = 22.767, df = 10, p-value = 0.01164
##
## Model df: 0. Total lags used: 10
##
##
## #### Residual checks for AMZN
## Series: train
## ARIMA(0,1,1) with drift
##
## Coefficients:
##             ma1    drift
##            -0.0327  0.0898
## s.e.      0.0205  0.0530
##
## sigma^2 = 7.427: log likelihood = -5987.39
## AIC=11980.77   AICc=11980.78   BIC=11998.21
```



```
##
## Ljung-Box test
##
## data: Residuals from ARIMA(0,1,1) with drift
## Q* = 7.0892, df = 9, p-value = 0.6278
##
## Model df: 1. Total lags used: 10
##
##
## 
## #### Residual checks for MSFT
## Series: train
## ARIMA(0,1,1) with drift
##
## Coefficients:
##          ma1    drift
##        -0.0811  0.1869
## s.e.   0.0205  0.0756
##
## sigma^2 = 16.73: log likelihood = -6991.42
## AIC=13988.85   AICc=13988.86   BIC=14006.29
```



```
##  
## Ljung-Box test  
##  
## data: Residuals from ARIMA(0,1,1) with drift  
## Q* = 19.156, df = 9, p-value = 0.0239  
##  
## Model df: 1. Total lags used: 10  
##  
##  
##  
##  
## #### Residual checks for TSLA  
## Series: train  
## ARIMA(0,1,0)  
##  
## sigma^2 = 54.33: log likelihood = -8448.99  
## AIC=16899.98 AICc=16899.99 BIC=16905.8
```



```
##  
## Ljung-Box test  
##  
## data: Residuals from ARIMA(0,1,0)  
## Q* = 23.389, df = 10, p-value = 0.009397  
##  
## Model df: 0. Total lags used: 10
```

Residual diagnostics evaluate whether systematic structure remains unexplained by the selected models. Autocorrelation patterns and Ljung–Box test results indicate that residuals for some securities deviate from white noise assumptions. This suggests potential benefits from alternative model classes or extensions incorporating additional dynamics.

Comparative Interpretation

Comparative Forecastability (Lower RMSE = More Predictable)

```
# Stock that is easiest/hardest to forecast (based on best RMSE)  
forecastability_tbl <- best_by_stock %>%  
  arrange(RMSE) %>%  
  mutate(forecastability_rank = row_number())  
  
forecastability_tbl %>%  
  knitr::kable(digits = 3,  
              caption = "Comparative Forecastability of Stocks")
```

Table 3: Comparative Forecastability of Stocks

ticker	model	RMSE	MAE	MAPE	forecastability_rank
AAPL	Naive	9.381	7.587	2.860	1
AMZN	Naive	21.271	19.656	8.518	2
TSLA	Naive	35.506	30.770	7.184	3
MSFT	Naive	39.808	35.896	7.602	4

Securities are ranked by minimum out-of-sample RMSE as a proxy for forecastability. Lower error values indicate more stable short-horizon dynamics under the evaluated modeling framework. The resulting ranking highlights heterogeneity in predictability across technology stocks.

Deployment

Reproducible Forecast Function

A simple deployment-ready function that can be reused in a script, Shiny app, or scheduled job.

```
forecast_stock <- function(ticker, from = start_date, to = end_date, h = 20){
  x <- getSymbols(ticker, src = "yahoo", from = from, to = to, auto.assign = FALSE)
  close <- as.numeric(Cl(x))

  fit <- auto.arima(close)
  fc  <- forecast(fit, h = h)

  list(
    ticker = ticker,
    model  = fit,
    forecast = fc
  )
}

# Example:
demo <- forecast_stock("AAPL", from = start_date, to = end_date, h = 20)
autoplot(demo$forecast) + labs(title = "Deployment Demo: AAPL 20-Day Forecast")
```



The forecasting workflow is encapsulated in a reusable function that retrieves recent market data, fits an ARIMA model, and generates short-term forecasts. This modular structure supports integration into dashboards, scheduled pipelines, or interactive applications, subject to appropriate production safeguards.

Conclusion

Key Findings

The analysis demonstrates meaningful differences in trend behavior, volatility regimes, and forecast accuracy across major technology stocks. Model effectiveness is asset-specific, and simple baselines remain competitive in several cases.

Next Improvement

Future work may incorporate alternative model families, exogenous predictors, extended forecast horizons, and automated deployment pipelines to enhance predictive performance and operational usability.

References

- Hyndman, R. J., & Athanasopoulos, G. (2018). *Forecasting: principles and practice*. OTexts.
- Yahoo Finance. (n.d.). Retrieved from <https://finance.yahoo.com/>
- R Documentation for `quantmod`, `forecast`, and `tseries` packages.
- Wickham, H., et al. (2019). *Welcome to the tidyverse*. Journal of Open Source Software, 4(43), 1686.
- Brockwell, P. J., & Davis, R. A. (2016). *Introduction to time series and forecasting*. Springer.
- Chatfield, C. (2003). *The analysis of time series: an introduction*. CRC press.
- Box, G. E., Jenkins, G. M., Reinsel, G. C., & Ljung, G. M. (2015). *Time series analysis: forecasting and control*. John Wiley & Sons.
- Tsay, R. S. (2010). *Analysis of financial time series*. John Wiley & Sons.
- Shumway, R. H., & Stoffer, D. S. (2017). *Time series analysis and its applications: with R examples*. Springer.
- ChatGPT (2024). *Assistance with R programming and time series analysis*.

Thank You!